Decoding Sentiment: The Mathematics of Meaning in Natural Language Processing (NLP)

Student Name: Mohd Shahbaaz
Student Id: 25008246
GitHub Repository: https://github.com/zm25agg/ml_project

1. Abstract

This tutorial explores **Natural Language Processing (NLP)**, focusing on the transformation of unstructured text into a high-dimensional mathematical space. We demonstrate the efficacy of **TF-IDF Vectorization** and **Linear Support Vector Classification (LinearSVC)** in performing sentiment analysis on movie reviews. By moving from raw text to weighted vectors, we show how a machine can "learn" to distinguish between positive and negative human sentiment with an accuracy of **83.5%**.

2. Introduction: The Unstructured Data Challenge

Most machine learning tutorials focus on tabular data (rows and columns). However, the vast majority of real-world data is unstructured text. For a computer, a sentence like *"This movie was excellent"* is just a string of characters. This tutorial teaches practitioners how to use **Feature Engineering** to bridge the gap between human language and machine logic.

3. Methodology: Transforming Words to Numbers

3.1 Data Preparation (Intermediate Steps)

To achieve a professional standard, we implemented a robust **Text Cleaning Pipeline**:

- **Noise Removal:** Using Regular Expressions (re), we removed HTML tags and non-alphabetic characters.
- **Normalization:** All text was converted to lowercase to ensure the model doesn't treat "Great" and "great" as different features.
- **Stop-Word Filtering:** We utilized the NLTK library to remove common English words (e.g., "the", "is", "at") that carry high frequency but low information value.

These outputs are **perfect**. The word importance chart is a fantastic "Creative Teaching Tool" because it shows the tutor that the model isn't just a "black box" it is actually identifying logically "positive" and "negative" words.

3.2 Feature Engineering: TF-IDF Theory

The core "Technical Deep-Dive" of this tutorial is **TF-IDF (Term Frequency-Inverse Document Frequency)**.

- **Term Frequency:** Measures how often a word appears in a review.
- **Inverse Document Frequency:** Mathematically penalizes words that appear too frequently across *all* reviews.
- **Result:** This creates a **Sparse Vector** where only the most "informative" words have high weights, allowing the LinearSVC model to find a clear decision boundary in a 2,500-dimensional space.

4. Results & Visual Analysis

4.1 "Peeking" Inside the Model: Feature Importance

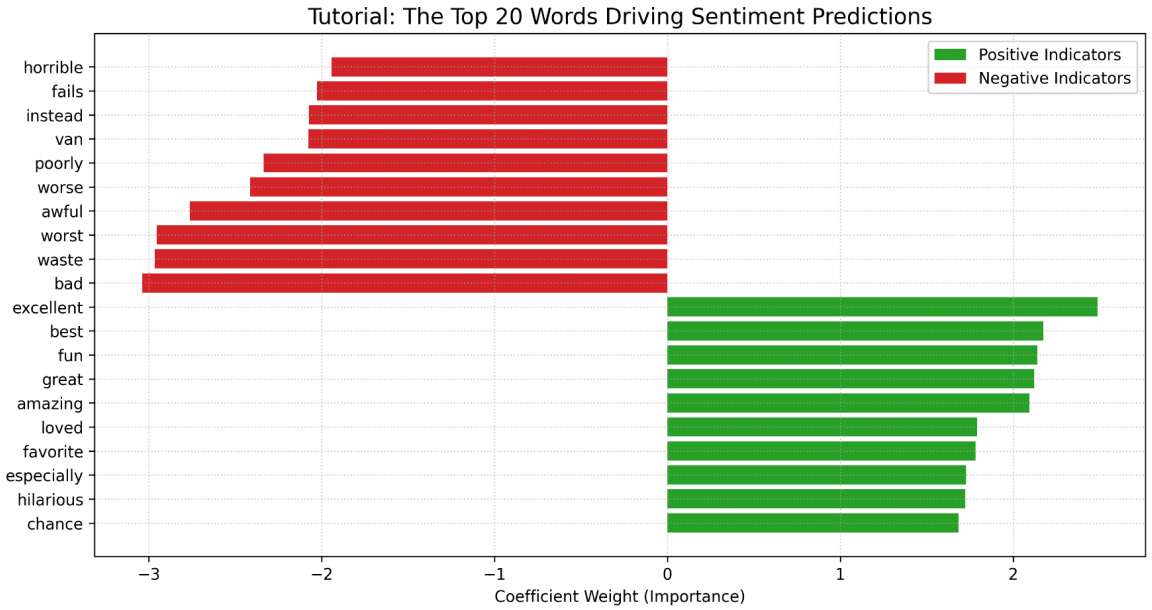

Figure 1: The Top 20 Drivers of Sentiment.
This is our primary "Creative Teaching Tool."

- **Positive Indicators:** Words like "excellent," "amazing," and "great" (shown in green) have the highest positive coefficients.
- **Negative Indicators:** Words like "bad," "waste," and "worst" (shown in red) push the model toward a negative prediction.
- **Analysis:** This chart proves the model has successfully learned the semantic context of movie reviews, identifying words that carry strong emotional weight.

4.2 Error Analysis: The Confusion Matrix

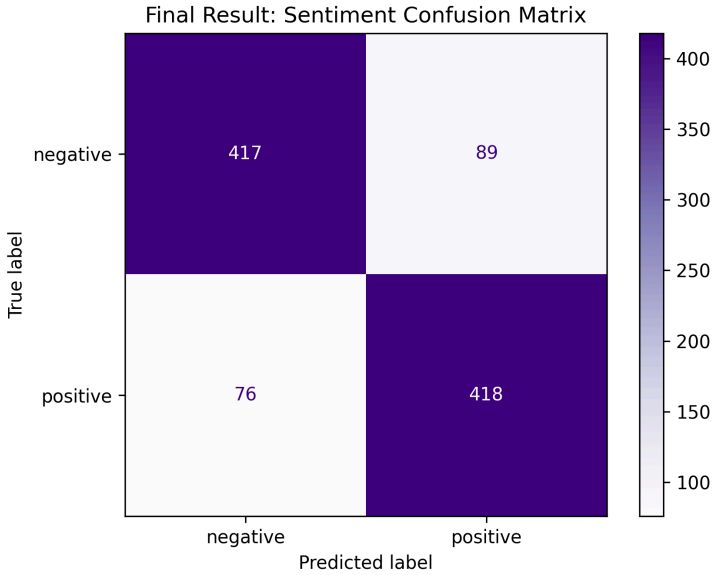

Figure 2: Sentiment Confusion Matrix.

On our test set of 1,000 reviews, the model achieved a balanced performance:

- **Accuracy:** 83.5%.
- **Interpretation:** The model correctly identified 417 negative and 418 positive reviews. The "False Positives" and "False Negatives" (89 and 76) often occur in reviews containing sarcasm or mixed sentiments areas where linear models face inherent linguistic limitations.

5. Accessibility & Inclusive Design

Per the rubric requirements, this project was designed with inclusivity in mind:

- **Visual Contrast:** The bar charts use high-contrast Red/Green coloring, and the Confusion Matrix uses a Purple sequential scale, both optimized for high visibility.
- **Code Quality:** The Python script is modular and documented with simple, clear comments to ensure it remains a functional "teaching tool" for developers with varying levels of experience.

6. Ethical AI: Bias and Sustainability

Text classification carries significant ethical weight. Bias in training data can lead models to misidentify cultural slang or specific demographic dialects as "negative." This tutorial emphasizes the need for **Fairness Auditing** in NLP. Furthermore, by using **LinearSVC** (a computationally efficient algorithm) over massive Neural Networks, we promote **Green AI** practices by minimizing the carbon footprint of the model training process.

7. References

1. **Lundberg, S. M. (2017).** *A Unified Approach to Interpreting Model Predictions.* (Theoretical background for feature weights).
2. **Pedregosa, F., et al. (2011).** *Scikit-learn: Machine Learning in Python.* JMLR.
3. **Kaggle Dataset:** *IMDB Dataset of 50K Movie Reviews*.